

# IP\_5th\_Random\_Forest\_ensemble\_methods

December 30, 2023

## 0.0.1 How To Learn Machine Learning Algorithms For Interviews

### Random Forest Classifier And Regresor

11. Ensemble Techniques(Boosting And Bagging)
12. Working of Random Forest Classifier
13. Working of Random Forest Regresor
14. Hyperparameter Tuning(Grid Search And RandomSearch)

### Xgboost Classifier And Regressor, GB Algorithm, Adaboost Decision Tree Theoretical Understanding:

1. Tutorial 37:Entropy In Decision Tree [https://www.youtube.com/watch?v=1IQOtJ4NI\\_0](https://www.youtube.com/watch?v=1IQOtJ4NI_0)
2. Tutorial 38:Information Gain <https://www.youtube.com/watch?v=FuTRucXB9rA>
3. Tutorial 39:Gini Impurity <https://www.youtube.com/watch?v=5aIFgrrTqOw>
4. Tutorial 40: Decision Tree For Numerical Features: <https://www.youtube.com/watch?v=5O8HvA9pMew>
5. How To Visualize DT: <https://www.youtube.com/watch?v=ot75kOmpYjI>

### Theoretical Understanding:

1. Ensemble technique(Bagging): <https://www.youtube.com/watch?v=KIOeZ5cFZ50>
2. Adaboost(Boosting Technique):<https://www.youtube.com/watch?v=NLRO1-jp5F8>
3. Gradient Boosting In Depth Intuition Part 1: <https://www.youtube.com/watch?v=Nol1hVtLOSg>
4. Gradient Boosting In Depth Intuition Part 2: <https://www.youtube.com/watch?v=Oo9q6YtGzvc>
5. Xgboost Classifier Indepth Intuition: <https://www.youtube.com/watch?v=gPciUPwWJQQ>
6. Xgboost Regression Indpeth Intuition: [https://www.youtube.com/watch?v=w\\_\\_vmVfpssg](https://www.youtube.com/watch?v=w__vmVfpssg)
7. Implementation of Xgboost: <https://youtu.be/9HomdnM12o4>

## 1 Random Forest

### Important properties of Random Forest Classifiers

1. Decision Tree—Low Bias And High Variance
2. Ensemble Bagging(Random Forest Classifier)—Low Bias And Low Variance

[ ]:

1. **What Are the Basic Assumption?** There are no such assumptions

[ ]:

## 2. Advantages Advantages of Random Forest

1. Doesn't Overfit
2. Favourite algorithm for Kaggle competition
3. Less Parameter Tuning required
4. Decision Tree can handle both continuous and categorical variables.
5. No feature scaling required: No feature scaling (standardization and normalization) required in case of Random Forest as it uses DEcision Tree internally
6. Suitable for any kind of ML problems

[ ]:

## 3. Disadvantages Disadvantages of Random Forest

1. Biased With features having many categories
2. Biased in multiclass classification problems towards more frequent classes.

[ ]:

## 4. Whether Feature Scaling is required? No

## 6. Impact of outliers? Robust to Outliers

[ ]:

## 4. Whether Feature Scaling is required? No

## 6. Impact of outliers? Robust to Outliers

[ ]:

## Types of Problems it can solve(Supervised)

1. Classification
2. Regression

[ ]:

## 2 Practical implementation:

<https://www.geeksforgeeks.org/random-forest-regression-in-python/>

```
[119]: import pandas as pd
import sklearn
import seaborn as sb
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
```

```
[120]: df=pd.read_csv("dataset/Employee_Salary_Dataset.csv")
```

```
[121]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 35 entries, 0 to 34
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                     35 non-null    int64
1   Experience_Years       35 non-null    int64
2   Age                   35 non-null    int64
3   Gender                 35 non-null    object
4   Salary                 35 non-null    int64
dtypes: int64(4), object(1)
memory usage: 1.5+ KB
```

```
[122]: df.shape
```

```
[122]: (35, 5)
```

```
[123]: df
```

```
[123]:
```

	ID	Experience_Years	Age	Gender	Salary
0	1	5	28	Female	250000
1	2	1	21	Male	50000
2	3	3	23	Female	170000
3	4	2	22	Male	25000
4	5	1	17	Male	10000
5	6	25	62	Male	5001000
6	7	19	54	Female	800000
7	8	2	21	Female	9000
8	9	10	36	Female	61500
9	10	15	54	Female	650000
10	11	4	26	Female	250000
11	12	6	29	Male	1400000
12	13	14	39	Male	6000050
13	14	11	40	Male	220100
14	15	2	23	Male	7500
15	16	4	27	Female	87000

16	17	10	34	Female	930000
17	18	15	54	Female	7900000
18	19	2	21	Male	15000
19	20	10	36	Male	330000
20	21	15	54	Male	6570000
21	22	4	26	Male	25000
22	23	5	29	Male	6845000
23	24	1	21	Female	6000
24	25	4	23	Female	8900
25	26	3	22	Female	20000
26	27	1	18	Male	3000
27	28	27	62	Female	10000000
28	29	19	54	Female	5000000
29	30	2	21	Female	6100
30	31	10	34	Male	80000
31	32	15	54	Male	900000
32	33	20	55	Female	1540000
33	34	19	53	Female	9300000
34	35	16	49	Male	7600000

```
[ ]:
```

```
[124]: l=LabelEncoder()
```

```
[125]: df["Gender"]=l.fit_transform(df["Gender"])
```

```
[126]: x=df.iloc[:,0:4].values
x
```

```
[126]: array([[ 1,  5, 28,  0],
               [ 2,  1, 21,  1],
               [ 3,  3, 23,  0],
               [ 4,  2, 22,  1],
               [ 5,  1, 17,  1],
               [ 6, 25, 62,  1],
               [ 7, 19, 54,  0],
               [ 8,  2, 21,  0],
               [ 9, 10, 36,  0],
               [10, 15, 54,  0],
               [11,  4, 26,  0],
               [12,  6, 29,  1],
               [13, 14, 39,  1],
               [14, 11, 40,  1],
               [15,  2, 23,  1],
               [16,  4, 27,  0],
               [17, 10, 34,  0],
               [18, 15, 54,  0],
```

```
[19,  2, 21,  1],
[20, 10, 36,  1],
[21, 15, 54,  1],
[22,  4, 26,  1],
[23,  5, 29,  1],
[24,  1, 21,  0],
[25,  4, 23,  0],
[26,  3, 22,  0],
[27,  1, 18,  1],
[28, 27, 62,  0],
[29, 19, 54,  0],
[30,  2, 21,  0],
[31, 10, 34,  1],
[32, 15, 54,  1],
[33, 20, 55,  0],
[34, 19, 53,  0],
[35, 16, 49,  1]], dtype=int64)
```

```
[127]: y=df.iloc[:, -1].values
y
```

```
[127]: array([ 250000,   50000,  170000,   25000,   10000, 5001000,
  800000,    9000,   61500,  650000,  250000, 1400000,
 6000050,  220100,    7500,   87000,  930000, 7900000,
  15000,  330000, 6570000,   25000, 6845000,    6000,
   8900,   20000,    3000, 10000000, 5000000,    6100,
  80000,  900000, 1540000,  9300000, 7600000], dtype=int64)
```

```
[128]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```

```
[129]: y_train
```

```
[129]: array([ 7600000, 6845000,   90000,  9300000, 1540000,   25000,
   25000,   80000,   65000, 10000000, 5001000,  220100,
 1400000,   15000,    9000,   250000, 5000000,   87000,
    3000, 6000050,   80000, 7900000,   10000,  930000,
   8900, 6570000], dtype=int64)
```

```
[130]: s=StandardScaler()
s.fit(x_train)
s.transform(x_train)
s.transform(x_test)
```

```
[130]: array([[ 0.34904982, -0.91792842, -0.87060341,  1.          ],
 [-0.33403693, -1.18019369, -1.07210806,  1.          ],
 [-1.6026266 , -1.31132632, -1.20644449,  1.          ],
 [ 1.12972039, -1.18019369, -1.20644449, -1.          ]],
```

```

[ 0.7393851 , -1.04906106, -1.13927628, -1.        ],
[ 0.54421746, -1.31132632, -1.20644449, -1.        ],
[-1.50504278, -1.04906106, -1.07210806, -1.        ],
[-0.91953985, -0.13113263, -0.19892126, -1.        ],
[ 0.15388218, -0.13113263, -0.19892126,  1.        ]])

```

```
[ ]:
```

```
[ ]:
```

```
[131]: re=RandomForestRegressor(oob_score=True)
```

```
[132]: re.fit(x_train,y_train)
```

```
[132]: RandomForestRegressor(oob_score=True)
```

```
[133]: re.oob_score_
```

```
[133]: 0.09668858029680272
```

```
[134]: from sklearn.metrics import mean_squared_error
```

```
[135]: y_pred=re.predict(x_test)
```

```
[136]: mean_squared_error(y_test,y_pred)
```

```
[136]: 976521079719.1389
```

```
[137]: re.score(x_test,y_test)
```

```
[137]: -92.0546149055722
```

```
[138]: #overfitted model
```

```
[ ]:
```

### 3 Ada boost

```
[139]: from sklearn.tree import DecisionTreeRegressor
       from sklearn.ensemble import AdaBoostRegressor
```

```
[140]: dt=DecisionTreeRegressor(max_depth=1)
```

```
[141]: ada=AdaBoostRegressor(dt, n_estimators=50, random_state=42)
```

```
[142]: ada.fit(x_train,y_train)
```

```
[142]: AdaBoostRegressor(estimator=DecisionTreeRegressor(max_depth=1), random_state=42)
```

```
[143]: # Make predictions on the test set
       predictions = ada.predict(x_test)

       # Evaluate Mean Squared Error
       mse = mean_squared_error(y_test, predictions)
       print("Mean Squared Error:", mse)
```

Mean Squared Error: 1792370368390.9731

```
[144]: ada.score(x_train,y_train)
```

```
[144]: 0.5001169877588378
```

```
[145]: ada.score(x_test,y_test)
```

```
[145]: -169.79849873466245
```

```
[ ]:
```

```
[ ]:
```

```
[146]: x_train
```

```
[146]: array([[35, 16, 49, 1],
           [23,  5, 29, 1],
           [32, 15, 54, 1],
           [34, 19, 53, 0],
           [33, 20, 55, 0],
           [ 1,  5, 28, 0],
           [ 4,  2, 22, 1],
           [ 7, 19, 54, 0],
           [10, 15, 54, 0],
           [28, 27, 62, 0],
           [ 6, 25, 62, 1],
           [14, 11, 40, 1],
           [12,  6, 29, 1],
           [19,  2, 21, 1],
           [ 8,  2, 21, 0],
           [11,  4, 26, 0],
           [29, 19, 54, 0],
           [16,  4, 27, 0],
           [27,  1, 18, 1],
           [13, 14, 39, 1],
           [31, 10, 34, 1],
           [18, 15, 54, 0],
           [ 5,  1, 17, 1],
           [17, 10, 34, 0],
           [25,  4, 23, 0],
```

```
[21, 15, 54, 1]], dtype=int64)
```

## 4 GBoost

```
[147]: from sklearn.ensemble import GradientBoostingRegressor
```

```
[148]: s=StandardScaler()  
s.fit(x_train)  
x_train=s.transform(x_train)  
x_test=s.transform(x_test)
```

```
[ ]:
```

```
[ ]:
```

```
[149]: # Create Gradient Boosting regressor  
gb_regressor = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1,  
↳max_depth=1, random_state=42)  
  
# Train the Gradient Boosting regressor  
gb_regressor.fit(x_train, y_train)  
  
# Make predictions on the test set  
predictions = gb_regressor.predict(x_test)  
  
# Evaluate Mean Squared Error  
mse = mean_squared_error(y_test, predictions)  
print("Mean Squared Error:", mse)
```

Mean Squared Error: 2120741249121.087

```
[150]: gb_regressor.score(x_train,y_train)
```

```
[150]: 0.7924476779635051
```

```
[151]: gb_regressor.score(x_test,y_test)
```

```
[151]: -201.08960600019404
```

## 5 i have to do feature engineering,hyper parameter tuning

```
[ ]:
```



## 6 XG Boost

```
[152]: import xgboost
```

```
[153]: from xgboost import XGBRegressor
```

```
[156]: # Create XGBoost regressor
xgb_regressor = XGBRegressor(n_estimators=100, learning_rate=0.1, max_depth=3,
    ↪random_state=42)

# Train the XGBoost regressor
xgb_regressor.fit(x_train, y_train)

# Make predictions on the test set
predictions = xgb_regressor.predict(x_test)

# Evaluate Mean Squared Error
mse = mean_squared_error(y_test, predictions)
print("Mean Squared Error:", mse)
```

Mean Squared Error: 976058570083.7365

```
[158]: xgb_regressor.score(x_train,y_train)
```

```
[158]: 0.99743131653229
```

```
[159]: xgb_regressor.score(x_test,y_test)
```

```
[159]: -92.01054145246775
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```